

Minor E. Gordon, Ph.D.

Curriculum Vitae

cv256@minorgordon.net

<http://mg4.consulting>

Tulsa, OK

Education

- 2009 **Ph.D., Computer Science**, University of Cambridge
Dissertation: Stage scheduling for CPU-intensive servers
- 2005 **Diplom Informatiker**, Technische Universität Berlin
Diplomarbeit: Staged design for highly concurrent web servers
- 2003 **B.Sc., Computer Science**, Oklahoma State University

Skills

Languages	English (native); German (fluent)
Programming	Java/J2EE; Python; C++; SQL; C#; TypeScript; PHP; JavaScript; HTML+CSS
Frameworks	Thrift; Protobufs; Guice; Guava; Vaadin; GWT; Shiro; jQuery; Bootstrap
Formats	XML; JSON; CSV; YAML; XDR
Testing	JUnit 4; Hamcrest; jMock; Jenkins; Pyunit; nose; gcov, gprof
Documentation	Doxygen; epydoc; Javadoc; Jsdoc
Databases	ElasticSearch; MongoDB; Redis; SQLite; H2; SQL Server
Platforms	Windows; Debian-, RedHat, and SuSE-based Linux; OS X
IDEs and tools	Visual C++; Eclipse CDT/JDT/PyDev; PyCharm; Maven; CMake; gcc, gmake, gdb
Revision control	Git; Mercurial; Subversion
Administration	Shell scripting; Nagios; Ganglia; Apache httpd and Tomcat; nginx

Commercial Projects

- 9/2014 – ongoing **Software architecture and backend development for the Notablist email newsletter search engine**
Part-time New York, NY USA and remote
- Notablist indexes millions of email newsletters from hundreds of thousands of publishers and provides both a search interface and real-time change alerts to users.
- Activities Prototyping in Python:
- Classified signup responses and mapped form inputs using numpy/scipy/scikit-learn
- Production rewrite in Java:
- Generated Python JSON-RPC clients, Python command line tools, Java service interfaces and abstract implementations, and Java JSON-RPC servlets from Thrift interface and data structure definitions
 - Implemented Guice-injected Java micro-services that encapsulated MongoDB collections; ElasticSearch indices; Redis databases; S3 buckets; screenshotting and signups with Selenium; SpamAssassin checking; logstash queries; Stripe integration; Drip (CRM) integration; MailChimp, SparkPost, and SendGrid transactional email posting; etcd locking; DNS and whois querying and parsing; DMOZ category and Alexa and Quantcast rank lookups
 - Designed and implemented a scalable distributed system for processing Common Crawl WARC records and submitting newsletter signups
 - Created administrative user interface in Vaadin
 - Docker+Compose deployment on real hardware

7/2012 – ongoing Full-stack development of the TeraScript web application server

Full-time then Tronics Software LLC / Stirling, NJ USA and remote

part-time TeraScribe and TeraScript Server are a visual development environment and server-based runtime for TeraScript Action Files (TAFs), TeraScript Class Files (TCFs), and the TeraScript Markup Language (TML). TML is a markup-based web application language, similar in spirit to ColdFusion. TAF and TCF are ways of organizing TML.

Activities TeraScript Server 8:

Rewrote TeraScript 7 in Java 7 and 8 while maintaining strict backwards compatibility:

- ANTLRv4-based compilers for TML and associated little languages
- Tree interpreters for TML, TAF, and TCF
- Versioned OSGi bundles with Maven and Apache Felix
- Library of standard functions, primitive types, and collections with extensive unit tests

Taught the product's owner the fundamentals and best practices of software design and implementation in Java so that he could maintain and improve the code base himself

TeraScript Server 7 (formerly Witango Server):

Reorganized, cleaned up, and modernized a 17-year-old C++ code base and addressed numerous bugs and feature requests for the server's first major release in over two years; now in maintenance releases

Eliminated diverging platform-specific build systems (on Win32, Linux, and OS X) in favor of CMake and ported the code base to Win64 using Visual C++ 2012

Designed and implemented a SQL generator that visits dialect-specific trees of SQL-99 constructs

TeraScribe 8:

Reorganized and cleaned up a 10-year-old Java Swing code base in order to add a number of features to the code editor, including autocomplete and syntax highlighting, and refit the data source management interface to use JDBC metadata

Subcontracting:

Cross-platform (Android, iOS, Windows) mobile app development with Xamarin Forms and C#

4/2017 – 5/2017 Drug name comparison project

Part-time PSW Applied Research Inc. / Toronto, ON and remote

Activities Implemented Kondrak's phonetic alignment and similarity algorithm (ALINE) and a bigram-based orthographic algorithm (BI-SIM) as a PostgreSQL C extension

Created a JavaScript demo application using jQuery and Bootstrap

4/2016 – 6/2016 Full-stack development for Polygon Analytics

Part-time Polygon Analytics Ltd. / London, UK and remote

Activities Implemented a geocoding and geotagging library for text documents, using libpostal and proprietary natural language processing algorithms in C++ (Visual Studio, CMake, Linux port)

Wrapped the library in Python with Boost.Python and as a REST API

Implemented a signup, payment (Stripe), and administration interface in Python with Flask and Redis

4/2014 – 7/2014 Application development for the Muninn Project under contract to the Library and Archives Canada

Part-time PSW Applied Research Inc. / Toronto, ON and remote

Activities Created a command-line Java application for transcribing handwritten World War I medical records from the Canadian Expeditionary Force (<http://blog.muninn-project.org/node/79>) using Amazon's Mechanical Turk crowdsourcing service: precut images for groups of lines, redundancy between workers, artificially-introduced mistakes and edit distances to detect cheating, and feedback to workers

3/2013 – 7/2013	Back end development for <u>Txtmrkt</u>, a marketplace for digital publications
Part-time	Txtmrkt LLC / New York, NY USA and remote
Activities	<p>Designed and implemented a Model-View-Controller-based Java web application with Guice-injected Thrift services backed by Amazon S3</p> <p>Authenticated and authorized users with Shiro</p> <p>Integrated co-branded order placement and payment notification with Amazon's Flexible Payments Service</p>
12/2011 – 6/2012	Haystack: performance monitoring platform for online businesses
Full-time	GrokIO LLC / New York, NY USA
Activities	<p>Designed an agent (Python) and manager (Java with Spring Security+MVC and Guice) architecture for gathering, storing, and querying time series data and metadata</p> <p>Implemented time series databases from scratch using relational tables and memory-mapped files</p> <p>Wrote RFC 3986 grammar-based URL and URN parsers using Ragel (Java) and pyparsing (Python)</p> <p>Generated service interfaces in Java and Python with custom Protocol Buffers compiler plugins</p> <p>Coded agent plugins in Java and Python for discovering and fetching metric metadata and data:</p> <ul style="list-style-type: none"> • Open source monitoring systems: Nagios, Ganglia, collectl, Graphite • Open source server software: Apache httpd, nginx, MySQL, memcached, MongoDB, JMX • Platform interfaces: procfs, sysctl, WMI, SNMP • Third party APIs: Google Analytics, MailChimp, Spring Social, Amazon CloudWatch <p>Supervised user experience and graphic designers in Ukraine</p> <p>Implemented a client-side user interface with the Google Web Toolkit, starting from wireframes:</p> <ul style="list-style-type: none"> • Architectural best practices: Model-View-Presenter; event bus; activities and places • Custom widgets: date time range picker, cell table pager, search box, selection tree • .war deployment to Tomcat with Maven
8/2011 – 12/2011	Back end development for <u>Birchbox</u>, a personalized subscription service for women's cosmetics samples
Full-time	Birchbox Inc. / New York, NY, USA
Activities	<p>Designed, implemented, documented, and successfully deployed warehouse integration code in Java (back end) and PHP (administrative front end)</p> <p>Wrote a Thrift compiler and targets in Python+pyparsing to generate Java and PHP from Thrift definitions</p> <p>Produced ad hoc business reports from MySQL using Java and Jython</p>
5/2011 – 8/2011	Benchmarking server stacks in order to accurately provision EC2 instances at <u>Chartbeat</u>, a real-time web analytics software-as-a-service provider
Full-time	Chartbeat Inc. / New York, NY, USA
Activities	<p>Synthesized a realistic benchmark for front-line web servers from statistical analyses of nginx access logs</p> <p>Significantly reduced network bandwidth and latency on production servers with targeted optimizations</p> <p>Wrote Python scripts for analyzing Ganglia RRDs, replaying HTTP request streams, harnessing httpfs</p>
10/2010 – 4/2011	File I/O library for <u>X10</u>, a type-safe, parallel object-oriented language for high-productivity computing
Full-time	IBM Research / Hawthorne, NY, USA
Activities	<p>Designed, implemented, tested, and documented a new low-overhead, buffer-based file library for X10:</p> <ul style="list-style-type: none"> • inspired by POSIX, FUSE, Boost.Filesystem, and Java NIO.2 • X10 and native code for the Java and C++ source-to-source compilation backends • Scatter/gather I/O, aligned buffers, memory-mapped files, advisory locking <p>Microbenchmarked the X10 runtime</p>

1/2009 – 7/2010 **LXFS: fast and reliable data storage for High-Performance Computing clusters**
 Full-time NEC High Performance Computing Europe / Stuttgart, Germany
 Activities Wrote Python and bash scripts for configuring, deploying, and administering LXFS installations:

- Lustre and Inet configuration (MGS, MDS, OSS, clients)
- Redundant NFS exports
- Nagios, Ganglia, and collectl monitoring
- Heartbeat/Linux-HA services for failover
- Promise RAID devices
- Network interfaces (Infiniband, bonded Ethernet)

Finished two major LXFS releases, used in numerous deployments
 Debugged and resolved issues in production parallel file systems

1/2009 – 7/2010 **XtreemFS: a distributed and replicated file system for wide-area networks**
 Full-time NEC High Performance Computing Europe / Stuttgart, Germany
 Activities Designed, implemented, and tested the XtreemFS userspace client:

- FUSE and Dokan (Windows FUSE-like library) interfaces
- Multiple pipelined ONC-RPC streams to a single server
- Staged, event-driven concurrency for robust performance under load
- Close-to-open file caching with per-file page sizes
- Automatic failover between file replicas on timeout
- Heavily benchmarked under different I/O loads (iobench, dbench, metadata benchmarks)

Finished two major XtreemFS releases, both used in production
 Collaborated with academic and industry partners in Europe, Israel, and China

Teaching Experience

2/2014 – 4/2014 Two sections of „SQL Fundamentals“ and „Python fundamentals“ workshops for 47 (total) Ph.D. students and post-doctoral researchers in the natural sciences
 Iowa State University / Ames, IA USA

2002 – 2004 Technical supervision of 10 Masters' theses in the U.S. and Poland
 Oklahoma State University / Tulsa, OK USA

1/2002 – 3/2002 Mathematics tutoring for 6th and 7th graders
 Tulsa Public Schools / Tulsa, OK USA